



기출 문제 & 답안

▼ 상태	CKA 자격증 학습
🕒 생성일시	@2022년 4월 1일 오전 7:41
👤 편집자	seongmi Lee
🕒 최종 편집일시	@2022년 4월 25일 오후 2:37

실전문제풀이

출판권 피카지 Online.

CKA 쿠버네티스 자격증 과정

PART1 | CKA 자격증 소개
자격증 등록과 Hands-On 실습 환경 만들기

PART2 | Kubernetes 아키텍처
Etcd 백업/복구, Kubernetes 업그레이드, RBAC 인증 실습

PART3 | Workloads & Scheduling
Application deploy 후 scale/rollback, pod 스케줄링 실습

PART4 | Services & Networking
Service 운영과 접근제한(NetworkPolicy), Ingress 운영 실습

PART5 | Storage
StorageClass 적용한 PV, PVC 생성하여 pod에 적용하기 실습

PART6 | Troubleshooting
controller component 설정 정보 수정 및 node/pod 문제해결

PART7 | 실전문제풀이
실전 문제를 직접 풀어 보고, 답안 리뷰 확인

1

? 1. Retrieve Error Messages from a Container Log

- Cluster: `kubectl config use-context hk8s`

In the `customa` namespace, check the log for the `nginx` container in the `custom-app` Pod.

Save the lines which contain the text “`error`” to the file `/var/CKA2022/errors.txt`.

▼ ! Answer

```
kubectl get pods -n customa
kubectl logs custom-app -n customa | grep -i error > /var/CKA2022/error.txt
cat /var/CKA2022/error.txt
```

? 2. Node Troubleshooting

- Cluster: `kubectl config use-context hk8s`

A Kubernetes worker node, named **hk8s-w2** is in state **NotReady**.

Investigate why this is the case, and perform any appropriate steps to bring the node to a **Ready** state, ensuring that any changes are made **permanent**.

▼ ! Answer

```
$ kubectl get nodes

$ ssh hk8s-w2
$ sudo -i
# docker ps
# systemctl status docker
# systemctl status kubelet
# systemctl enable --now kubelet
# systemctl status kubelet
# exit
$ exit

$ kubectl get nodes
```

? 3. Count the Number of Nodes That Are Ready to Run Normal Workloads

- Cluster: kubectl config use-context **hk8s**

Determine how many nodes in the cluster are **ready** to run normal workloads (i.e., workloads that do not have any special tolerations).

Output this number to the file **/var/CKA2022/count.txt**

▼ ! Answer

```
kubectl get nodes | grep -i -w ready | wc -l > /var/CKA2022/count.txt
cat /var/CKA2022/count.txt
3
```

? 4. Management Node

- Cluster: kubectl config use-context **k8s**

Set the node named **k8s-worker1** as **unavailable** and **reschedule all the pods running on it**.

▼ ! Answer

```
kubectl drain k8s-worker2 --ignore-daemonsets --force

kubectl get nodes

kubectl get pods
```

? 5. ETCD backup & restore

- 작업 클러스터 : kubectl config use-context **k8s**

First, create a snapshot of the existing etcd instance running at <https://127.0.0.1:2379>, saving the snapshot to **/data/etcd-snapshot.db**.

Next, restore an existing, previous snapshot located at **/data/etcd-snapshot-previous.db**.

The following TLS certificates/key are supplied for connecting to the server with etcdctl:

CA certificate: **/etc/kubernetes/pki/etcd/ca.crt**

Client certificate: `/etc/kubernetes/pki/etcd/server.crt`
Client key: `/etc/kubernetes/pki/etcd/server.key`

▼ ! Answer

```
ssh k8s-master
sudo -i
sudo ETCDCTL_API=3 etcdctl --endpoints=https://127.0.0.1:2379 \
  --cacert=/etc/kubernetes/pki/etcd/ca.crt \
  --cert=/etc/kubernetes/pki/etcd/server.crt \
  --key=/etc/kubernetes/pki/etcd/server.key \
  snapshot save /data/etcd-snapshot.db

sudo ETCDCTL_API=3 etcdctl --data-dir=/var/lib/etcd-previous snapshot restore /data/etcd-snapshot-previous.db
sudo tree /var/lib/etcd-previous/

sudo vi /etc/kubernetes/manifests/etcd.yaml
...
- hostPath:
  path: /var/lib/etcd-previous
  type: DirectoryOrCreate
  name: etcd-data

sudo docker ps -a | grep etcd
exit
exit
```

? 6. Cluster Upgrade - only Master

- 작업 클러스터 : `kubectl config use-context k8s`

upgrade system : `hk8s-m`

Given an existing Kubernetes cluster running version `1.22.4`,

upgrade all of the Kubernetes control plane and node components on the master node only to version `1.23.3`.

Be sure to `drain` the `master` node before upgrading it and `uncordon` it after the upgrade.

▼ ! Answer

```
# kubeadm 업그레이드
sudo yum install -y kubeadm-1.23.3-0 --disableexcludes=kubernetes
kubeadm version

# node components 업그레이드
sudo kubeadm upgrade plan v1.23.3
sudo kubeadm upgrade apply v1.23.3

# 노드 드레인
kubectl drain hk8s-m --ignore-daemonsets

# kubelet과 kubectl 업그레이드
sudo yum install -y kubelet-1.23.3-0 kubectl-1.23.3-0 --disableexcludes=kubernetes
sudo systemctl daemon-reload
sudo systemctl restart kubelet

# 노드 uncordon
sudo kubectl uncordon hk8s-m
```

? 7. Authentication and Authorization

- Cluster : `k8s`

Context You have been asked to create a new `ClusterRole` for a deployment pipeline and bind it to a specific `ServiceAccount` scoped to a specific namespace.

Task:

- Create a new `ClusterRole` named `deployment-clusterrole`, which only **allows to create** the following resource types:

Deployment StatefulSet DaemonSet

- Create a new **ServiceAccount** named `cicd-token` in the existing namespace `app-team1`.
- Bind the new **ClusterRole** `deployment-clusterrole` to the new **ServiceAccount** `cicd-token`, limited to the namespace `app-team1`.

▼ ! Answer

```
kubectl create clusterrole deployment-clusterrole --verb=create --resource=deployment,statefulset,daemonset
kubectl get clusterrole deployment-clusterrole

kubectl create serviceaccount cicd-token --namespace=app-team1
kubectl get serviceaccounts --namespace app-team1

kubectl create clusterrolebinding deployment-clusterrolebinding --clusterrole=deployment-clusterrole --serviceaccount=app-team1:cicd-token
kubectl describe clusterrolebindings deployment-clusterrolebinding
```

? 8. Pod 생성하기

- 작업 클러스터 : `kubectl config use-context k8s`

Create a **new namespace** and create a **pod** in the namespace

TASK:

- namespace name: `cka-exam`
- pod Name: `pod-01`
- image: `busybox`
- environment Variable: `CERT = "CKA-cert"`
- command: `/bin/sh`
- args: `-c "while true; do echo $(CERT); sleep 10;done"`

▼ ! Answer

```
kubectl create namespace cka-exam
kubectl get namespaces cka-exam

kubectl run pod-01 --image=busybox --dry-run=client -o yaml > pod-01.yaml
vi pod-01.yaml
apiVersion: v1
kind: Pod
metadata:
  name: pod-01
  namespace: cka-exam
spec:
  containers:
  - env:
    - name: CERT
      value: CKA-cert
    image: busybox
    name: pod-01
    command: ["/bin/sh"]
    args: ["-c", "while true; do echo $(CERT); sleep 10;done"]

kubectl apply -f pod-01.yaml
kubectl get pod -n cka-exam
```

? 9. multi-container Pod 생성

- cluster : `kubectl config use-context hk8s`
- Create a pod with 4 containers running : `nginx`, `redis`, `memcached` and `consul`
 - pod name: `eshop-frontend`
 - image: `nginx`
 - image: `redis`

- image: memcached
- image: consul

▼ ! Answer

```
kubectl run eshop-frontend --image=nginx --dry-run=client -o yaml > multi.yaml

vi multi.yaml
apiVersion: v1
kind: Pod
metadata:
  name: eshop-frontend
spec:
  containers:
  - image: nginx
    name: nginx
  - image: redis
    name: redis
  - image: memcached
    name: memcached
  - image: consul
    name: consul

kubectl apply -f multi.yaml
```

? 10. Side-car Container Pod 실행

- 작업 클러스터 : `kubectl config use-context k8s`
- 현재 운영 중인 **eshop-cart-app** Pod의 로그를 Kubernetes built-in logging 아키텍처(예: `kubectl logs`)에 통합하는 로그 스트리밍 사이드카 컨테이너를 운영하시오.
 - busybox 이미지를 사용하여 **price** 라는 이름의 side-car container를 기존 **eshop-cart-app**에 추가합니다.
 - 새 price 컨테이너는 다음과 같은 command를 실행해야 합니다.
Command: `/bin/sh, -c, "tail -n+1 -f /var/log/cart-app.log"`
 - /var/log에 마운트 된 볼륨을 사용하여 사이드카 컨테이너에서 로그 파일 **cart-app.log**를 사용해야 합니다.
 - **eshop-cart-app Pod와 cart-app 컨테이너를 수정하지 마시오**

▼ ! Answer

```
kubectl get pod eshop-cart-app -o yaml > sidecar.yaml
cat sidecar.yaml
apiVersion: v1
kind: Pod
metadata:
  name: eshop-cart-app
spec:
  containers:
  - image: busybox
    name: cart-app
    command: ['/bin/sh', '-c', 'i=1;while :;do echo -e "$i: Price: $((RANDOM % 10000 + 1))" >> /var/log/cart-app.log; i=$((i+1));']
    volumeMounts:
    - name: varlog
      mountPath: /var/log
  volumes:
  - emptyDir: {}
    name: varlog

vi sidecar.yaml
apiVersion: v1
kind: Pod
metadata:
  name: eshop-cart-app
spec:
  containers:
  - image: busybox
    name: cart-app
    command: ['/bin/sh', '-c', 'i=1;while :;do echo -e "$i: Price: $((RANDOM % 10000 + 1))" >> /var/log/cart-app.log; i=$((i+1));']
    volumeMounts:
    - name: varlog
      mountPath: /var/log
  - name: price
```

```

    image: busybox
    args: [/bin/sh, -c, "tail -n+1 -f /var/log/cart-app.log"]
    volumeMounts:
      - name: varlog
        mountPath: /var/log
  volumes:
    - emptyDir: {}
      name: varlog

kubectl delete pod eshop-cart-app
kubectl apply -f multi.yaml

```

? 11. Pod Scale-out

- Cluster: kubectl config use-context **k8s**

Expand the number of running Pods in "**eshop-order**" to 5.

- namespace: **devops**
- deployment: **eshop-order**
- replicas: **5**

▼ ! Answer

```

kubectl deployment -n devops
kubectl get deployment eshop-order -n devops
kubectl scale deployment eshop-order --replicas=5 --namespace devops

```

? 12. Rolling Update

- Cluster: kubectl config use-context **k8s**

Create a deployment as follows:

- TASK:
 - name: **nginx-app**
 - Using container **nginx** with version **1.11.10-alpine**
 - The deployment should contain **3** replicas
- Next, deploy the application with **new version 1.11.13-alpine**, by performing a **rolling update**
- Finally, **rollback** that update to the **previous version 1.11.10-alpine**

▼ ! Answer

```

# create
kubectl create deployment nginx-app --image=nginx:1.11.10-alpine --replicas=3 --dry-run=client -o yaml > deployment.yaml
cat deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: nginx-app
    name: nginx-app
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx-app
  strategy: {}
  template:
    metadata:

```

```

    creationTimestamp: null
    labels:
      app: nginx-app
  spec:
    containers:
      - image: nginx:1.11.10-alpine
        name: nginx
        resources: {}
  status: {}

kubectrl apply -f deployment.yaml --record

# rolling update
kubectrl set image deployment nginx-app nginx=nginx:1.11.13-alpine --record

kubectrl rollout history deployment nginx-app

# roll-back
kubectrl rollout undo deployment nginx-app
kubectrl rollout history deployment nginx-app

```

? 13. Network Policy with Namespace

- 작업 클러스터 : kubectrl config use-context **k8s**

Create a new **NetworkPolicy** named **allow-port-from-namespace** in the existing namespace **devops**.

Ensure that the new NetworkPolicy allows Pods in namespace **migops**(using label **team=migops**) to connect to port **80** of Pods in namespace **devops**.

Further ensure that the new NetworkPolicy: does not allow access to Pods, which don't listen on **port 80** does not allow access from Pods, which are not in namespace **migops**

▼ ! Answer

```

vi policy.yaml
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: allow-port-from-namespace
  namespace: devops
spec:
  podSelector: {}
  policyTypes:
    - Ingress
  ingress:
    - from:
      - namespaceSelector:
          matchLabels:
            team: migops
      ports:
        - protocol: TCP
          port: 80

kubectrl apply -f policy.yaml

```

? 14. Create a persistent volume

- Cluster: kubectrl config use-context **k8s**

Create a persistent volume with name **app-config**, of capacity **1Gi** and access mode **ReadWriteMany**.

The type of volume is **hostPath** and its location is **/var/app-config**.

▼ ! Answer

```

vi pv.yaml
apiVersion: v1
kind: PersistentVolume

```

```

metadata:
  name: app-config
spec:
  capacity:
    storage: 16i
  accessModes:
    - ReadWriteMany
  hostPath:
    path: /var/app-config

```

```

kubectl apply -f pv.yaml
kubectl get pv

```

? 15. Deploy and Service

- 작업 클러스터 : kubectl config use-context k8s

Reconfigure the existing deployment `front-end` and add a port specification named `http` exposing port `80/tcp` of the existing container `nginx`.

Create a new service named `front-end-svc` exposing the container port `http`.

Configure the new service to also expose the individual Pods via a `NodePort` on the nodes on which they are scheduled

▼ ! Answer

```

kubectl get deploy front-end -o yaml > front-end.yaml

cat > front-end.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: front-end
spec:
  selector:
    matchLabels:
      app: front-end
  replicas: 2
  template:
    metadata:
      labels:
        app: front-end
    spec:
      containers:
        - name: http
          image: nginx
          ports:
            - name: http
              containerPort: 80
---
apiVersion: v1
kind: Service
metadata:
  name: front-end-svc
spec:
  type: NodePort
  ports:
    - port: 80
      protocol: TCP
      targetPort: http
  selector:
    app: front-end

kubectl delete deploy front-end
kubectl apply -f front-end.yaml
kubectl get deployments.apps,svc

```

? 16. DNS Lookup

- 작업 클러스터 : kubectl config use-context k8s

Create a `nginx` pod called `nginx-resolver` using image `nginx`, expose it internally with a service called `nginx-resolver-service`.

Test that you are able to look up the service and pod names from within the cluster. Use the image: `busybox:1.28` for `dns`

lookup.

- Record results in `/var/CKA2022/nginx.svc` and `/var/CKA2022/nginx.pod`
- Pod: `nginx-resolver` created
- Service DNS Resolution recorded correctly
- Pod DNS resolution recorded correctly

▼ ! Answer

```
#Create a nginx pod
kubectl run nginx-resolver --image=nginx --port=80
kubectl expose pod nginx-resolver --name nginx-resolver-service --port=80
kubectl get svc nginx-resolver-service
  nginx-resolver-service  10.104.150.11
# Test:Service DNS Resolution recorded correctly -> /var/CKA2022/nginx.svc
kubectl run test-nslookup --image=busybox:1.28 -it --restart=Never --rm -- nslookup 10.104.150.11
kubectl run test-nslookup --image=busybox:1.28 -it --restart=Never --rm -- nslookup nginx-resolver-service > /var/CKA2022/nginx.s

# Test: Pod DNS resolution recorded correctly -> /var/CKA2022/nginx.pod
kubectl get pod nginx-resolver -o wide
  nginx-resolver  10.244.1.55
kubectl run test-nslookup --image=busybox:1.28 -it --restart=Never --rm -- nslookup 10-244-1-55.default.pod.cluster.local > /var/

$ cat /var/CKA2022/nginx.svc
$ cat /var/CKA2022/nginx.pod
```

? 17. Application with PVC

- Cluster: kubectl config use-context `k8s`

Create a new PersistentVolumeClaim:

Name: `pv-volume`

Class: `csi-hostpath-sc`

Capacity: `10Mi`

Create a new Pod which mounts the PersistentVolumeClaim as a volume:

Name: `web-server`

Image: `nginx`

Mount path: `/usr/share/nginx/html`

Configure the new Pod to have `ReadWriteOnce` access on the volume.

▼ ! Answer

```
vi pvc.yaml
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: pv-volume
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 10Mi
  storageClassName: csi-hostpath-sc

kubectl apply -f pvc.yaml
kubectl get pv,pvc

vi pod-with-pvc.yaml
apiVersion: v1
kind: Pod
metadata:
  name: web-server
spec:
  containers:
    - name: web-server
      image: nginx
```

```
    volumeMounts:
      - mountPath: "/usr/share/nginx/html"
        name: pv-volume
  volumes:
    - name: pv-volume
      persistentVolumeClaim:
        claimName: pv-volume

kubectl apply -f pod-with-pvc
kubectl get pod
```